

Customer Churn Prediction System using Python, Machine Learning, EDA, SMOTE, and Deep Learning

Project Report

Abstract

Customer churn is one of the biggest challenges faced by companies across industries such as telecommunications, banking, insurance, and e-commerce. Losing existing customers not only reduces revenue but also increases the cost of acquiring new customers. This project focuses on building a Customer Churn Prediction System using Python, Machine Learning, Exploratory Data Analysis (EDA), SMOTE, and Deep Learning to identify customers who are likely to leave a company.

The project follows a complete machine learning pipeline, starting with data preprocessing and visualization, followed by feature engineering, handling class imbalance using SMOTE, training multiple machine learning models, and developing an Artificial Neural Network (ANN) for churn prediction. The results demonstrate how data-driven techniques can assist businesses in making informed decisions and implementing customer retention strategies.

Chapter 1: Introduction

Customer churn refers to the situation where customers stop using a company's products or services. In highly competitive markets, customer retention is often more profitable than acquiring new customers. Businesses that can identify customers with a high probability of leaving can take preventive measures by offering discounts, personalized services, or improved customer support.

Machine Learning has become an effective solution for predicting customer churn by analyzing customer behavior and identifying hidden patterns within historical data. This project combines traditional machine learning algorithms with deep learning techniques to develop an accurate prediction system.

The project also emphasizes proper data analysis using Exploratory Data Analysis (EDA), preprocessing techniques, and SMOTE (Synthetic Minority Oversampling Technique) to improve model performance when dealing with imbalanced datasets.

Chapter 2: Objectives

The primary objectives of this project are:

- To analyze customer data using Exploratory Data Analysis.
- To preprocess the dataset by handling missing values and categorical features.
- To identify important features affecting customer churn.
- To solve the class imbalance problem using SMOTE.
- To build machine learning models for churn prediction.
- To develop a Deep Learning model using an Artificial Neural Network.
- To evaluate model performance using standard classification metrics.
- To assist businesses in improving customer retention strategies.

Chapter 3: Technologies Used

Programming Language:

- Python

Libraries:

- Pandas
- NumPy
- Matplotlib
- Seaborn
- Scikit-learn
- TensorFlow
- Keras
- imbalanced-learn (SMOTE)

Development Environment:

- Google Colab

Chapter 4: Methodology

4.1 Data Collection

The project begins with collecting customer information containing demographic details, subscription information, account history, service usage, and customer status. The target variable indicates whether a customer has churned.

4.2 Data Preprocessing

Before training the models, the dataset is cleaned and prepared by:

- Removing duplicate records
- Handling missing values
- Encoding categorical variables
- Feature scaling
- Splitting the dataset into training and testing sets

Data preprocessing improves the quality of the dataset and ensures that machine learning models can learn meaningful patterns.

4.3 Exploratory Data Analysis (EDA)

Exploratory Data Analysis helps understand customer behavior and relationships among variables.

Several visualizations are created to analyze:

- Customer distribution
- Churn rate
- Contract type
- Payment methods
- Monthly charges
- Total charges
- Customer tenure
- Correlation between features

EDA provides valuable business insights and helps identify factors that influence customer churn.

4.4 Handling Imbalanced Data using SMOTE

Most customer churn datasets contain significantly fewer churn cases than non-churn cases. Such imbalance causes machine learning models to become biased toward the majority class.

To overcome this issue, SMOTE (Synthetic Minority Oversampling Technique) is applied.

SMOTE generates synthetic samples for the minority class instead of simply duplicating existing records. This creates a balanced dataset and allows the models to learn both classes more effectively.

The use of SMOTE improves model fairness and increases recall for churn prediction.

4.5 Machine Learning Models

Multiple supervised machine learning algorithms can be trained and compared.

Examples include:

- Logistic Regression
- Decision Tree
- Random Forest
- Support Vector Machine
- K-Nearest Neighbors

Each model learns patterns from historical customer data and predicts whether a customer is likely to churn.

4.6 Deep Learning Model

In addition to traditional machine learning models, an Artificial Neural Network (ANN) is developed.

The neural network consists of:

- Input Layer
- Hidden Layers
- Activation Functions
- Output Layer

The ANN captures complex nonlinear relationships between customer attributes and churn behavior.

TensorFlow and Keras are used to implement the deep learning model.

Chapter 5: Model Evaluation

The performance of the models is evaluated using several classification metrics.

Accuracy

Measures the percentage of correctly classified customers.

Precision

Measures how many predicted churn customers actually churned.

Recall

Measures how many actual churn customers were correctly identified.

F1-Score

Balances precision and recall into a single performance metric.

Confusion Matrix

Shows the number of correct and incorrect predictions made by the model.

ROC-AUC Score

Measures the model's ability to distinguish between churn and non-churn customers.

Using multiple evaluation metrics provides a better understanding of model performance than relying only on accuracy.

Chapter 6: Results and Discussion

The project successfully demonstrates the complete workflow of customer churn prediction using data science and machine learning techniques.

Exploratory Data Analysis helped uncover important customer trends and relationships.

Data preprocessing improved the overall quality of the dataset.

Applying SMOTE effectively reduced the impact of class imbalance and improved prediction capability for minority-class customers.

Machine learning models provided reliable classification performance, while the Artificial Neural Network demonstrated the ability to learn more complex patterns within the data.

The comparison of different approaches highlights the importance of selecting appropriate preprocessing techniques and evaluation metrics for real-world predictive analytics.

Chapter 7: Applications

Customer churn prediction has numerous real-world applications, including:

- Telecommunications
- Banking
- Insurance
- E-commerce
- Healthcare
- Online subscription services
- Financial institutions
- SaaS companies

Businesses can use churn prediction systems to identify at-risk customers and implement targeted retention strategies.

Chapter 8: Advantages

- Improves customer retention.
- Reduces customer acquisition costs.

- Supports business decision-making.
- Automates customer risk prediction.
- Handles imbalanced datasets effectively.
- Combines Machine Learning and Deep Learning techniques.
- Provides meaningful business insights through data visualization.

Chapter 9: Limitations

- Model performance depends on data quality.
- Predictions may vary with different datasets.
- Customer behavior changes over time, requiring periodic retraining.
- Deep learning models require greater computational resources.

Several enhancements can further improve this project.

Future improvements include:

- Hyperparameter optimization.
- Deployment using Flask, FastAPI, or Streamlit.
- Integration with cloud platforms.
- Explainable AI using SHAP or LIME.
- Real-time customer churn prediction.
- Automated model retraining pipelines.
- Integration with CRM systems.

Conclusion

This project presents a complete Customer Churn Prediction System developed using Python, Machine Learning, Exploratory Data Analysis, SMOTE, and Deep Learning. The workflow includes data preprocessing, visualization, feature engineering, handling imbalanced data, model development, and performance evaluation.

The project demonstrates how machine learning techniques can be applied to solve real-world business problems. By predicting customers who are likely to churn, organizations can design effective retention strategies, improve customer satisfaction, and increase long-term profitability.

Overall, the project provides practical experience in data analysis, predictive modeling, and deep learning while highlighting the importance of data-driven decision-making in modern businesses.

References

1. Scikit-learn Documentation.
2. TensorFlow Documentation.
3. Pandas Documentation.
4. NumPy Documentation.
5. Matplotlib Documentation.
6. Seaborn Documentation.
7. imbalanced-learn Documentation (SMOTE).
8. Research literature on customer churn prediction using machine learning and deep learning.

Chapter 10: Future Scope

The Customer Churn Prediction System can be further enhanced in several ways to improve its accuracy, scalability, and practical use in business environments.

- **Real-Time Churn Prediction:** Integrate the model with live customer databases or CRM systems to predict churn as customer data is updated.
- **Advanced Deep Learning Models:** Explore more sophisticated neural network architectures and ensemble learning techniques to improve predictive performance.
- **Hyperparameter Optimization:** Apply automated tuning methods such as Grid Search, Random Search, or Bayesian Optimization to identify the best model parameters.
- **Model Explainability:** Incorporate explainable AI tools such as SHAP or LIME to help businesses understand the factors influencing churn predictions and build trust in the model.
- **Web Application Deployment:** Deploy the model as a user-friendly web application using Flask, FastAPI, Streamlit, or Django, allowing businesses to make predictions without requiring programming knowledge.
- **Cloud Integration:** Host the application on cloud platforms such as AWS, Microsoft Azure, or Google Cloud to improve scalability, availability, and accessibility.

- **Automated Model Retraining:** Develop an automated pipeline that periodically retrains the model with new customer data, ensuring it remains accurate as customer behavior changes.
- **Support for Multiple Industries:** Adapt the system for use in banking, insurance, healthcare, e-commerce, subscription services, and other industries by training it on domain-specific datasets.
- **Enhanced Feature Engineering:** Incorporate additional customer information, such as transaction history, customer feedback, social media interactions, and service usage patterns, to improve prediction accuracy.
- **Business Decision Support:** Integrate the prediction system with customer retention strategies by automatically recommending personalized offers, discounts, or loyalty programs for customers identified as being at high risk of churning.

Presented by

Luis Gallo

Prepared Using:

- Python
- Machine Learning
- Deep Learning
- Exploratory Data Analysis (EDA)
- SMOTE
- Google Colab

Date: July 2026